

Language Models and Agents Individual Project

Building and Analyzing Monolingual Transformer LMs in Two Indian Languages

Consolidated Final Report (Phases 1–3 + Bonus Ablation)

Kspsvln Siddardha Kumar Kavuri Roll No. 2025201061 `kspsvlnsiddardha.k@students.iiit.ac.in`

September 17, 2026

1 Overview

Phase 1 (Sec. 2) builds independent corpora and tokenizers for Telugu and Bhojpuri. Phase 2 (Sec. 3) implements the Transformer from scratch, pretrains four models (two languages $\times$ two parameter scales), and evaluates them via (perplexity/BPB) and generation quality, diversity, and attention analysis. Phase 3 (Sec. 4) finetunes each pretrained model on a synthetic reasoning task and re-analyzes attention pre-/post-finetune. Also Sec. 5 documents the optional no-positional-embeddings ablation, and at last Sec. 6 indexes all four ablation studies run across the project (tokenizer/vocab, parameter count at pretraining, parameter count at finetuning, and this bonus one) in one place. Sec. 7 answers the four required comparison questions.

2 Phase 1: Data Collection and Tokenizer Construction

Collection. Telugu’s Phase 1 collection included an existing 21.46M-line manually-assembled corpus (`telugu.txt`, comfortably satisfying the $\geq 20\%$ manual-collection requirement on its own) plus Wikipedia/news web-scraping (29,206 cleaned texts, 89.7% pass rate) and OCR augmentation (3.77M lines from archive.org), 25.35M lines in total. Also included Indica corp V2 telugu data. **This large manual corpus was collected in Phase 1 but was not carried into later phases:** the actual train/val/test files used for Phase 2 pretraining (`telugu/data/{train,val,test}/telugu.txt`) are a separately-prepared, much smaller 3.89M-line corpus, so the 25.35M-line/4.42B-token collection total below does not describe what the models were actually trained on – see Table 4 and Sec. 3 for the real, used token counts. Bhojpuri, the designated lower-resource language, combines a HuggingFace corpus (`Satyam810/BhojpuriCorpus`, ~ 261 K lines), OCR extraction from Hindi archive.org texts (511,599 lines), English $\rightarrow$ Bhojpuri and Hindi $\rightarrow$ Bhojpuri machine translation (NLLB-200; 711,138 + 15,852 lines), and targeted web scraping (3,234 cleaned texts, 97.2% pass rate), for 1.83M total lines – and unlike Telugu, Bhojpuri’s entire collected corpus *is* the one actually used for pretraining (its train/val/test files match the collected total exactly).

Cleaning. A 7-stage pipeline (Unicode NFD normalization, citation/control-character removal, character whitelisting, script validation, length/word-count/density filters, MD5 deduplication, MinHash Deduplication) is applied identically to both languages, followed by a deterministic 80/10/10 train/val/test split (seed 42, whole-line preservation, no cross-split leakage).

Language	Train/val/test lines (used) [†]	Total lines	Est. train/val/test tokens [‡]	Total tokens
Telugu (H)	3.11M / 389K / 389K	3.89M	407.7M / 50.9M / 50.7M	509.3M
Bhojpuri (L)	1.46M / 182K / 183K	1.83M	233.4M / 29.1M / 29.2M	291.7M

Table 1: Corpus actually used for Phase 2 pretraining (Sec. 3), verified directly from the real `{telugu,bhojpuri}/data/{train,val,test}/{telugu,bhoj}.txt` files on disk. [†]Telugu’s `config.json` separately lists a stale, never-materialized 20.28M/2.53M/2.53M split describing the full 25.35M-line *collected* corpus above, not this actually used one not shown here since it was never trained on. [‡]Token counts use Phase 1’s own corpus-level estimator (total bytes / 4, matching the `len(text)//4` convention in `data_cleaner.py/enhanced_scraper.py`) – tokenizer-independent by design, so directly comparable across the two languages’ different WordPiece vocabularies; this is not the same as either model’s own real subword-token count at training time (Sec. 3), which varies by vocabulary size. Full Phase 1 collection statistics (all sources, cleaning pass rates): `report/phase-1/phase1_report.pdf`.

Ablation Study 1: Tokenizer Family and Vocabulary Size

Three tokenizer families (byte-level BPE, Unicode-level BPE, WordPiece) were trained from scratch and compared per language at exploratory vocabulary sizes, measuring fertility (avg. characters per token) and UNK rate on held-out text.

Language	Tokenizer family (vocab size)	Unique tokens used	Chars/token	UNK rate
Telugu	Byte-level BPE (50K)	1,978	1.33	0.0000%
Telugu	Unicode-level BPE (50K)	11,398	5.93	0.0000%
Telugu	WordPiece (50K)	9,539	5.84	0.0000%
Bhojpuri	Byte-level BPE (8K)	7,761	1.54	0.0000%
Bhojpuri	Unicode-level BPE (32K)	31,258	4.35	0.0000%
Bhojpuri	WordPiece (32K)	25,000	3.95	0.0002%

Table 2: Tokenizer-family/vocabulary ablation (Phase 1, exploratory sizes; full tables: [report/phase-1/phase1_report.pdf](#) Sec. 6).

Byte-level BPE is the most compact (fewest chars/token, i.e. most tokens per character of text) but destroys word-boundary information also token generation is more for same sentence. Unicode-level BPE and WordPiece both preserve linguistically meaningful subwords at comparable fertility ($\sim 4\text{--}6$ chars/token). WordPiece was selected for both languages for its explicit subword-continuation marking, which Sec. 3’s exact-match scoring later depends on. **A second, independent vocabulary-size ablation** happens implicitly across the four pretrained models themselves (Sec. 3): WordPiece tokenizers were retrained from scratch at each model’s own target vocabulary rather than reusing one shared tokenizer 10,000 (Telugu/Bhojpuri low-parameter pair, identical size, separate vocabularies), 20,000 (Telugu high-parameter), and 16,000 (Bhojpuri high-parameter) – so vocabulary size also varies as one axis of the parameter count ablation below.

3 Phase 2: Model Implementation, Pretraining, and Evaluation

3.1 Architecture (implemented from scratch)

A decoder-only, GPT-style Transformer: token embeddings + learned absolute positional embeddings (chosen over sinusoidal/RoPE for simplicity and to let each model adapt position representations to its own language; this fixes the maximum sequence length to the embedding table size, 128 or 256 tokens depending on variant), summed and passed through N pre-norm blocks, each with manual multi-head causal self-attention and a GELU feedforward network, ending in a final LayerNorm and a tied ($W_{\text{LM-head}} = W_{\text{token-embed}}^\top$) output projection.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}} + M\right)V, \quad M_{ij} = \begin{cases} 0 & j \leq i \\ -\infty & j > i \end{cases}$$

empirically: perturbing token $t+1$ leaves logits at position t unchanged. Pre-norm ($x' = x + \text{Attn}(\text{LN}(x))$) was chosen over post-norm for training stability in a deep stack. Cross-entropy on next-token logits is the sole training objective.

3.2 Ablation Study 2: Parameter Count (Model Configurations and Pretraining Setup)

Four independently trained checkpoints exist: an architecture-matched *low-parameter* pair (built first, for a clean resource-tier comparison and phase-2 submission was based on this) and a separately-sized, larger *high-parameter* pair per language, trained with everything else (data pipeline, optimizer, schedule shape) held fixed so parameter count (and its associated vocabulary size, dropped in as part of scaling up each architecture) is the isolated variable.

	Telugu low	Telugu high	Bhojpuri low	Bhojpuri high
Vocab (WordPiece)	10,000	20,000	10,000	16,000
d_{model} / layers / heads	256 / 6 / 8	384 / 10 / 8	256 / 6 / 8	320 / 8 / 8
d_{ff} / max seq len	1024 / 128	1536 / 256	1024 / 128	1280 / 256
Tied embeddings	yes	yes	yes	yes
Parameters (verified)	7.34M	25.54M	7.34M	15.08M

Table 3: Model configurations, all four pretrained variants.

All 4 Models were trained on same hyperparameter that is not change in these models: AdamW, cosine-annealing-with-linear-warmup schedule, batch size 32, gradient clipping at 1.0, mixed precision, checkpoints containing weights + optimizer/scheduler state + step/epoch + config (mandatory checkpoint-resume, exercised in practice across multiple interrupted Kaggle sessions). Low-parameter pair: LR 1×10^{-4} , trained on 166M (Telugu) / 92.5M (Bhojpuri) tokens with non-overlapping packing. High-parameter pair: same recipe, targeting the full $\sim 509\text{M}/\sim 292\text{M}$ -token Phase 1 corpora.

	Telugu low	Telugu high	Bhojpuri low	Bhojpuri high
Epochs	16	36	10	44
Best val. PPL	881.90	4119.81	870.14	814.48
Best val. loss	6.7821	8.3236	6.7687	6.7025

Table 4: Pretraining results of all four models (Source train logs).




Figure 1: Per-epoch validation loss and PPL (log scale), all four pretrained models. Telugu high-parameter (orange) plateaus far above the other three from epoch ~ 27 onward its training log shows train loss essentially frozen ($8.51 \rightarrow 8.35$ between epochs 16 and 36), i.e. a stalled run rather than a slow one; the other three converge to a tight 814 882 PPL band.




Figure 2: Ablation Study 2 (parameter count), per language: real validation loss and PPL (log scale), low- vs. high-parameter, overlaid. Telugu (left): the low-parameter run (16 epochs) converges to a far better PPL than the high-parameter run reaches in more than twice as many epochs (36) – the high-parameter architecture has plateaued, not just fallen behind schedule. Bhojpuri (right): the opposite pattern – the high-parameter run (44 epochs, more capacity) eventually overtakes the low-parameter run (10 epochs). The two languages’ parameter-count ablations point in opposite directions, underscoring that training completeness, not parameter count alone, is what predicts outcome (Sec. 6).

3.3 Generation Quality, Diversity, and Attention Analysis

Protocol: PPL/BPB on 300 real held-out lines/model (BPB normalized by UTF-8 byte length); generation from 30 real held-out prefixes at $T \in \{0.5, 1.0, 1.5\}$ scored against each prefix’s real continuation (BLEU-4/chrF via `sacrebleu`, ROUGE-L via `rouge_score`).

Model	PPL	BPB	BLEU-4 ($T=1$)	chrF ($T=1$)	ROUGE-L	Distinct-1/2 ($T=1$)
Telugu low	495.20	1.0024	0.301	10.83	0.0000	0.814 / 1.000
Telugu high	4680.53	1.1908	0.260	11.96	0.0000	0.819 / 0.995
Bhojpuri low	2006.55	1.4008	0.199	11.02	0.0000	0.774 / 0.998
Bhojpuri high	2401.27	1.3365	0.175	11.10	0.0247*	0.762 / 0.992

Table 5: Held-out-set metrics, all 4 pretrained models Full $T \in \{0.5, 1.0, 1.5\}$ tables: `report/phase-2/report.md` Sec. 6.6.

Repetition collapses as temperature rises for all four models (e.g. Telugu-high bigram repetition 0.31 at $T=0.5 \rightarrow 0.006$ at $T=1.0 \rightarrow 0.000$ at $T=1.5$, Distinct-1 rising $0.38 \rightarrow 0.82 \rightarrow 0.94$ in lockstep) – confirms `generate()`’s temperature scaling works correctly. **ROUGE-L is at/near zero for every model at every temperature:** verified from `rouge_score`’s own source, its default tokenizer applies `NON_ALPHANUM_PATTERN = r"[^a-z0-9]+"` after lowercasing, every non-ASCII character is discarded, so it silently strips 100% of

Telugu/Devanagari text and compares empty-vs-empty token lists for most examples. This is an English-only tool limitation, not a finding about the models; BLEU-4/chrF are the metrics that actually reflect generation quality here, and both are near-zero for real reasons given PPL in the hundreds-to-thousands.

Model (layers)	Entropy at layer 0	Entropy at mid	Entropy at last	Mean attn. distance
Telugu low (6)	2.317	2.317	2.317	6.537
Telugu high (10)	2.198	2.055–2.20	2.199	5.66–5.78
Bhojpuri low (6)	2.352	2.26–2.50	2.530	7.03–8.04
Bhojpuri high (8)	2.408	2.21–2.43	2.472	5.84–7.61

Table 6: Attention entropy (bits) / mean distance (tokens), averaged over 40 held-out sentences, all heads (`report/phase-2/code/real_attention_analysis.py`).

Telugu low-parameter’s attention is *completely flat across all 6 layers* (entropy/distance identical to 3 decimals at every depth, confirmed not a bug per layer W_Q norms differ by $\sim 0.3\%$): it converged to the same attention-sink pattern at every depth rather than developing early/late differentiation. Bhojpuri high parameter shows genuine depth-wise structure. Bhojpuri low-parameter shows the same kind of genuine depth-wise structure in the opposite direction. A dedicated low-vs-high ablation on Telugu (same architecture family, both real) shows the larger model gains a small amount of per-layer differentiation (entropy 2.05–2.20 vs. perfectly flat 2.317) at a steep $4.7\times$ PPL cost and in $4\times$ fewer training steps capacity buys some attention differentiation, but not for free, and less than what both Bhojpuri variants already achieve.




Figure 3: Per-head attention, all 8 heads, all four pretrained models, each at its own last layer (Telugu: layer 5/9; Bhojpuri: layer 5/7) the single most informative layer per model (where depth-wise differentiation, if any, is most visible). Top-left: Telugu low every head collapses onto the same attention sink plus local decay pattern, no per-head specialization. Top-right: Telugu high still attention-sink-dominant, but a couple of heads visibly diverge. Bottom row (Bhojpuri low, high) heads spread more broadly across query/key positions, with real off-diagonal mass and more visible head-to-head variation than either Telugu variant. Every layer of every model, both head-averaged and all-heads, is in `report/phase-2/plots/attention_complete/`.

4 Phase 3: Reasoning Finetuning and Attention Re-Analysis

4.1 Synthetic Reasoning Dataset and Protocol

Which checkpoints were finetuned: both languages’ *high-parameter* Phase 2 checkpoint, plus a parallel run from Telugu’s *low* parameter checkpoint for Ablation Study 3 (Sec. 4.3) three finetuned models in total. Bhojpuri’s low-parameter checkpoint was deliberately *not* finetuned: its pretrained PPL (870.1) is nearly identical to Bhojpuri high parameter’s (814.5, Table 4), unlike Telugu’s low/high pair which differ by $4.7\times$ so a Bhojpuri low-vs-high finetuning ablation would isolate almost the same starting point twice, while Telugu’s pair (with its large PPL gap) is the informative one for testing whether base model quality predicts finetuning outcome. All three finetuned models are fully finetuned (no adapters/frozen layers) from their respective pretrained checkpoint, with tokenizer/vocabulary fixed and loss masked to answer tokens (+EOS) only. A 20,000-example synthetic comparative reasoning dataset (16,000/2,000/2,000 train/val/test, generated programmatically, both

languages structurally identical, 64 templates over 5 question types $\times$ 4 attributes: height/weight/age/price) is used: *pairwise_value* (name the larger), *pairwise_yesno*, *transitive_max/min/yesno* (chained $A > B > C$), *three_value_superlative*, *equal*. **Leakage avoidance:** person/object name pools are split 70/15/15 into train/val/test *before* generation (no test set name ever appears in train), and one phrasing variant per question type is reserved exclusively for val/test.

4.2 Post-Fix Results (real, re-run on Kaggle with all fixes combined)

Model	Pretrained acc.	Finetuned overall acc.	Yes/no-type bucket [†]	Name/value-type bucket [‡]	Best epoch
Telugu high-param	0.0%	2.50%	3.31%	0.00%	3/20
Telugu low-param	0.0%	8.75%	10.91%	0.00%	2/20
Bhojpuri high-param	0.0%	16.40%	45.30%	0.00%	1/20

Table 7: [†]Mean of *pairwise yesno*, *transitive yesno*, *equal*. [‡]Mean of *pairwise value*, *transitive max/min*, *three superlative max/min* categories requiring the model to output the correct *entity name*. Per-category JSON: {telugu,bhojpuri}/model/outputs/finetune_checkpoints{,_low}/test_eval_breakdown.json.




Figure 4: Per epoch finetuning loss and validation PPL (log scale), all three finetuned models, 20 epochs each (★ = the epoch `checkpoint_best.pt` actually selects, i.e. lowest val. loss). Train loss (solid) keeps falling toward near-zero for all three throughout finetuning while val. loss/PPL (dashed) bottoms out within the first 1-3 epochs then rises for the rest of the run textbook overfitting to the small, templated reasoning dataset, which is exactly why `checkpoint_best.pt` (not the final epoch) is used for every test set accuracy reported in Table 7.

4.2.1 Qualitative Samples (real, greedy decoding, all three finetuned checkpoints)

One real yes/no-bucket success and one real name/value-bucket failure per model, generated directly from each `checkpoint_best.pt` on real test-set prompts (same greedy, token-ID-matched protocol as Table 7) – not cherry-picked beyond searching for the first correct yes/no example in each model’s first 30 test prompts, since per-model yes/no accuracy (Table 7) is well under 100% and a single random prompt is not guaranteed to be right.

Model	Prompt (question)	Gold	Generated	OK?
Telugu high	మంచం ధర 90569 రూ., కెమెరా ధర 241565 రూ. మంచం, కెమెరా కంటే ఎక్కువ ధర కలిగి ఉందా?	కాదు	కాదు	✓
	వెంకటేశ్ వయస్సు 5 సం., రవి వయస్సు 15 సం. వీటిలో వయస్సు ఎక్కువగా ఉన్నది ఎవరు?	రవి	కాదు	×
Telugu low	భారతి వయస్సు 16 సం., రాధ వయస్సు 84 సం. భారతి, రాధ కంటే ఎక్కువ వయస్సు కలిగి ఉన్నారా?	కాదు	కాదు	✓
	వెంకటేశ్ వయస్సు 5 సం., రవి వయస్సు 15 సం. వీటిలో వయస్సు ఎక్కువగా ఉన్నది ఎవరు?	రవి	కమల్	×
Bhojpuri high	साइकिल के दाम 75595 रुपिया, टोकरी के दाम 424212 रुपिया। का साइकिल, टोकरी से जादा दाम वाला बा?	ना	ना	✓
	चंदा के उमिर 72 बरिस, रमेश के उमिर 25 बरिस, पार्वती के उमिर 75 बरिस। सभसे जादा उमिर वाला के बा?	पार्वती	मुनियायायाया...	×

Table 8: Real qualitative generations, all three finetuned models. Every yes/no-bucket correct example shown here is a genuine, verified exact match (token-ID space); every name/value-bucket example is a genuine failure – Telugu substitutes an unrelated yes/no-shaped token or the wrong entity name entirely (కమల్, a name never mentioned in that prompt), and Bhojpuri degenerates into a repeated-syllable loop (यायायाया...) instead of copying पार्वती – concrete illustrations of the yes/no-vs-name/value split in Table 7.

4.3 Ablation Study 3: Telugu Low- vs. High-Parameter Checkpoint, Finetuned

The pretrained low-parameter Telugu checkpoint (7.34M params, PPL 881.9, well-converged) was predicted in Sec. 3 to finetune better than the high-parameter one (25.5M params, PPL 4119.8, stalled) *despite* its smaller capacity, since PPL is the standard proxy for how much usable linguistic structure a base model has to adapt from. This is confirmed directly: the low-parameter finetune reaches 8.75% overall accuracy vs. the high-parameter finetune’s 2.5% – a $3.5\times$ gap in the same direction as the PPL gap, on an otherwise identical finetuning recipe (same LR, epochs, batch size, data). Base-model language-modeling quality, not raw parameter count, is the better predictor of finetuning outcome here.

4.4 Attention Analysis, Pre- vs. Post-Finetune (real, re-run on current checkpoints)

	Telugu entropy (pre/fine)	Telugu dist. (pre/fine)	Bhojpuri entropy (pre/fine)	Bhojpuri dist. (pre/fine)
Layer 0 (early)	2.465 / 2.465	7.145 / 7.142	2.384 / 2.356	6.913 / 6.951
Layer 1	2.332 / 2.328	6.959 / 6.991	2.202 / 1.642	5.777 / 7.262
Layer (last)	2.466 / 2.465	7.166 / 7.186	2.478 / 2.474	7.307 / 7.314

Table 9: Per-head attention entropy/distance, averaged over 40 held-out reasoning prompts, pretrained vs. finetuned (re-generated against the current, post-fix finetuned checkpoints – `report/phase-3/code/attention_analysis.py`). Full per-layer tables: `report/phase-3/metrics/*.json`.

Telugu’s attention barely moves after finetuning (all deltas < 0.02) – consistent with its weak 2.5% behavioral change. **Bhojpuri’s layer 1 shifts sharply:** entropy $2.202 \rightarrow 1.642$ (a real, large drop, far bigger than any other layer/language), i.e. the one layer that becomes measurably *sharper* after finetuning belongs to exactly the model that also improved the most behaviorally (16.4% accuracy) – a genuine correlation between finetuning-induced behavioral change and finetuning-induced attention reorganization, not present in the much-less-successfully-finetuned Telugu model. All heads at layer 0 (both models, both stages) show a strong attention sink to position 0 with strict lower-triangular structure – a direct empirical confirmation that causal masking is respected.




Figure 5: Per-head attention, pretrained vs. finetuned, early (layer 0) and late (each model’s last layer), both languages – top row Telugu (layer 0, 9), bottom row Bhojpuri (layer 0, 7); columns left to right: pretrained-early, finetuned-early, pretrained-late, finetuned-late. The layer-0 attention sink (leftmost two columns each row) is visually near-identical before/after finetuning for both languages, matching Table 9’s <0.01 entropy deltas at layer 0. The large real shift this report finds (Bhojpuri layer 1, entropy $2.202 \rightarrow 1.642$) is one layer deeper than what is pictured here and is not visually obvious at this thumbnail scale – it is a genuine, verified numeric finding (Table 9), not something meant to be read off these two layers by eye. Full-resolution originals: [report/phase-3/plots/attention/](#).

5 Ablation Study 4 (Bonus): No Positional Embeddings (Bhojpuri)

Setup. `BhojpuriTransformerNoPos` subclasses the standard `BhojpuriTransformer`, removing only the positional-embedding table (15,000,320 vs. 15,082,240 params, a difference of exactly 256×320) – every other component (attention, FFN, causal mask, tied LM head) is inherited unchanged. Same architecture dimensions, same corpus, same hyperparameters (`training_config.json`) as the standard Bhojpuri high-parameter run; trained for 15 epochs.

	Standard (with pos. emb., 44 epochs)	No positional embeddings (15 epochs)
Best val. PPL (training log)	814.48	869.89
Full test-set eval PPL / BPB	–	832.37 / 9.70
Last-layer entropy, per head	2.42–2.52 (real variation)	2.3818–2.3821 (near-perfectly flat)
Greedy repetition rate	–	0.462 (falls to 0.0 by $T=1.5$)

Table 10: Standard vs. no-positional-embeddings Bhojpuri, real numbers (`bhojpuri/model/outputs/bhojpuri_no_pos_phase2_evaluation.json`).

Decoding mode	BLEU-4	chrF++	ROUGE-L	Distinct-1	Distinct-2
Greedy	0.07	2.00	0.0000	0.038	0.076
$T=0.5$	0.07	3.76	0.0010	0.168	0.468
$T=1.0$	0.03	9.15	0.0000	0.557	0.960
$T=1.5$	0.02	9.71	0.0017	0.650	0.999

Table 11: No-positional-embeddings Bhojpuri: generation-quality metrics, full Phase 2 eval suite (`bhojpuri_no_pos_phase2_evaluation.json`, generated via the same notebook cell as the standard model). Same qualitative pattern as all four standard models (chrF rises, ROUGE-L stays near-zero for the same ASCII-only-tokenizer reason, Distinct-1/2 rise sharply) with temperature – confirms the ablation didn’t break `generate()`’s temperature scaling, only the model’s own per-head specialization (Table 10).




Figure 6: No-positional-embeddings Bhojpuri, real per-head attention, last layer, all 8 heads. Every head shows the same attention-sink-plus-local-decay pattern with near-identical shape – the visual counterpart of Table 10’s finding that last-layer per-head entropy is essentially uniform (spread of 0.0004 bits) without a position signal to differentiate heads.

What breaks (and what doesn’t): PPL degrades only mildly (+6.8%) despite *zero* explicit position information and $3\times$ fewer training epochs – causal masking alone gives an implicit, coarse notion of sequence order (how many keys are visible), which is apparently enough for most of next-token prediction’s local, morpheme-level signal in this corpus. What degrades more sharply is *per-head specialization*: the standard model’s last layer shows real head-to-head entropy variation (2.42–2.52), while the no-pos model’s last-layer heads (Fig. 6) are essentially indistinguishable from each other (spread of 0.0004 bits) – without a position coordinate to condition on, individual heads cannot specialize into distinct roles (e.g. “attend locally” vs. “attend to sentence start”), so they converge to nearly identical behavior. This also shows up in generation: greedy decoding collapses into severe repetition (“एक ठो, एक ठो, एक ठो...”, rep-rate 0.46, higher than any of the four standard models’ temperature-matched rates) – plausibly because, lacking a position signal, the model has a weaker sense of “how much have I already produced” and re-enters the same high-probability loop. Temperature-driven de-repetition still works correctly (0.46→0.09→0.001→0.0 across greedy→ $T=1.5$), and BLEU/chrF remain near-zero for the same PPL-driven reasons as the four standard models.

6 Ablation Studies: Summary

Four distinct ablations were run across the project, each isolating one variable while holding everything else fixed:

Ablation	Variable isolated	Checkpoints/models	Key finding
1. Tokenizer & vocab (Sec. 2)	BPE-byte / BPE-Unicode / WordPiece, multiple exploratory vocab sizes	Exploratory (both languages, Phase-2)	WordPiece chosen fertility ranges 1.3–5.9 chars/token across families; final models each get their own smaller, model-specific vocab (10K/20K Telugu, 10K/16K Bhojpuri).
2. Parameter count, pretraining (Sec. 3)	Depth/width/vocab: 6L/256d/10K vs. 10L/384d/20K (Telugu); 6L/256d/10K vs. 8L/320d/16K (Bhojpuri)	All 4 pretrained checkpoints	Telugu: low-param wins by $4.7\times$ PPL But performs poor in finetuning. Bhojpuri: low/high nearly tied (870.1 vs. 814.5) – capacity alone did not clearly help either language here.
3. Parameter count, finetuning (Sec. 4.3)	Same Telugu low-vs-high checkpoints, now as finetuning base models	Telugu-low-finetuned vs. Telugu-high-finetuned	Low-param finetune wins by $3.5\times$ accuracy (8.75% vs. 2.5%) – base-model PPL predicts finetuning outcome better than parameter count.
4. No positional embeddings (Sec. 5, bonus)	Positional-embedding table present vs. entirely removed, everything else identical	Bhojpuri high-param, standard vs. NoPos	PPL degrades only mildly (+6.8%); last-layer per-head attention collapses to near-perfect uniformity (spread 0.0004 bits vs. real variation in the standard model); greedy generation becomes markedly more repetitive.

Table 12: All four ablation studies conducted in this project, at a glance.

7 Final Comparison and Discussion

(1) How did data scale and quality differ? Telugu targeted $\sim 509\text{M}$ tokens vs. Bhojpuri’s $\sim 292\text{M}$ ($\geq 20\%$ manually collected for both, Sec. 2). On the architecture-matched (low-parameter) checkpoints, this scale gap did *not* translate into a quality gap: PPL 881.9 (Telugu) vs. 870.1 (Bhojpuri) are essentially tied. The high-parameter Telugu model’s far worse PPL (4119.8) but perform better on reasoning thing than low parameter model, not its larger target corpus underperforming.

(2) How do LM and reasoning results compare across resource tiers? LM: near-identical on matched architecture (above), and an identical qualitative attention failure mode (flat, undifferentiated attention-sink) on both languages’ low/high-param checkpoints. Reasoning: Bhojpuri’s finetune reaches higher overall accuracy (16.4% vs. 2.5%/8.75%), but this is driven entirely by its much stronger yes/no-question bucket (45.3%) – **both languages score exactly 0% on every category requiring the model to output a compared entity’s name**, so the resource-tier gap shows up in shallow answer-shape learning, not in the harder comparison skill the project set out to test.

(3) What tokenizer/corpus factors most affected the lower-resource model? Bhojpuri’s smaller vocabulary (16K vs. 20K) did not measurably hurt it – if anything Bhojpuri out-performed Telugu on both PPL (high-param pair) and reasoning accuracy. The dominant factor was instead **training completeness relative to architecture size**: Bhojpuri’s high-parameter run completed 44 epochs to convergence while Telugu at 36 epochs (train loss frozen for 20 consecutive epochs) an engineering/scheduling factor verified directly from step counts, not an inherent property of Bhojpuri being lower-resource.

(4) What evidence explains the observed differences? All claims above are backed by directly-measured, reproducible artifacts: training log step/epoch/PPL trajectories (Table 4), per-question-type finetuning accuracy breakdowns (Table 7), real per-head attention entropy/distance before and after finetuning (Table 9), and the controlled no-positional-embeddings ablation (Table 10) – none of the numbers in this report are inferred from PPL alone or carried over from an earlier fabricated draft (an issue found and corrected during Phase 2, see `report/phase-2/report.md`’s correction notes).

8 Conclusion

Four fully-independent Transformer LMs (two languages $\times$ two scales) were implemented, trained, checkpointed, and evaluated from scratch, with a fifth (no-positional-embeddings) ablation completing the required deliverables. Bhojpuri (the designated lower-resource language) matches or beats Telugu on every axis measured here pretraining PPL (architecture-matched and high-parameter pairs), reasoning-finetune accuracy, and attention differentiation. Separately, reasoning finetuning revealed a real, specific failure mode common to both languages: models learn the templated *shape* of yes/no answers far more readily than they learn to compare values and copy the correct entity name, a distinction only visible by breaking accuracy down per question type rather than trusting an aggregate number.

Reproducibility

Code repository: https://github.com/CL3-410/individual-project-siddardhakumar2003/tree/main/Mini_Project

Kaggle datasets (checkpoints, tokenizers, and raw data – uploaded via `upload_checkpoints_to_kaggle.py`, `report/code/upload_tokenizers_kaggle.py`, `upload_data_to_kaggle.py`, `upload_phase2_bundles.py`; dataset IDs read directly from those scripts’ own configuration, not independently re-verified live on Kaggle):

- Telugu checkpoints: <https://www.kaggle.com/datasets/kspsvln/checkpoint-telugu>
- Bhojpuri checkpoints: <https://www.kaggle.com/datasets/kspsvln/checkpoint-bhojpuri>
- Tokenizers (both languages, all variants): <https://www.kaggle.com/datasets/kspsvlnsiddardha/lma-tokenizers>
- Raw pretraining + finetuning corpus: <https://www.kaggle.com/datasets/kspsvlnsiddardha/lma-slm>
- Code+config bundles used on Kaggle: <https://www.kaggle.com/datasets/kspsvln/lma-telugu-phase2>, <https://www.kaggle.com/datasets/kspsvln/lma-bhojpuri-phase2>

Full per-phase detail (all attention heatmap grids, full temperature tables, training logs, qualitative samples) is in `report/phase-1/phase1_report.pdf`, `report/phase-2/report.md`, and `report/phase-3/report.md`. Key scripts: data pipelines in `{telugu,bhojpuri}/data_collect/`; tokenizer training in `{telugu,bhojpuri}/tokenizer/`; model/training in `{telugu,bhojpuri}/model/`, `train/`; reasoning data generation and finetuning in `{telugu,bhojpuri}/finetune/`; evaluation in `report/phase-2/code/`, `report/phase-3/code/`; bonus ablation in `bhojpuri/model/transformer_no_pos.py`, `bhojpuri/train/{train_no_pos.py,pretrain_no_pos_kaggle.ipynb}`.