

Language Models and Agents (LMA)

Phase 1 Report

Data Collection & Preprocessing for Telugu and Bhojpuri

Kpsvln Siddardha Kumar Kavuri

Roll Number 2025201061

kpsvlnsiddardha.k@students.iiit.ac.in

August 19, 2026

Phase 1 Completion Date: August 19, 2026

Dataset Available on Kaggle

<https://www.kaggle.com/datasets/kpsvlnsiddardha/lma-slm>

Contents

1	Phase 1 Data Report	4
1.1	Data Collection Summary	4
1.1.1	Telugu (Model H)	4
1.1.2	Bhojpuri (Model L)	4
1.2	Data Processing Pipeline	4
1.2.1	Cleaning Process	4
1.2.2	Data Splits (80/10/10)	5
2	Data Collection Overview	5
2.1	Methodology	5
2.1.1	Telugu (Model H) - Higher-Resource Language	5
2.1.2	Bhojpuri (Model L) - Lower-Resource Language	5
2.2	Data Sources	5
3	Data Cleaning Pipeline	6
3.1	7-Stage Cleaning Process	6
3.2	Quality Metrics	6
3.2.1	Telugu Cleaning Analysis	6
3.2.2	Bhojpuri Cleaning Analysis	6
4	Data Statistics & Visualizations	7
4.1	Collection Metrics Comparison	7
4.2	Cleaning Pipeline Results	7
4.3	Train/Validation/Test Split Distribution	8
4.4	Data Source Composition (Bhojpuri)	8

5	Train/Validation/Test Splits	8
5.1	Split Strategy	8
5.2	Telugu Splits	9
5.3	Bhojpuri Splits	9
6	Detailed Statistics	9
6.1	Telugu (Model H) Statistics	9
6.2	Bhojpuri (Model L) Statistics	10
7	Tokenizer Training	10
7.1	Telugu Tokenizer (Model H) - Three Variants	10
7.1.1	Variant 1: Byte-Level BPE (50K Vocab)	10
7.1.2	Variant 2: Unicode-Level BPE (50K Vocab)	11
7.1.3	Variant 3: WordPiece (50K Vocab)	11
7.2	Bhojpuri Tokenizer (Model L) - Three Variants	11
7.2.1	Variant 1: Byte-Level BPE (8K Vocab)	11
7.2.2	Variant 2: Unicode-Level BPE (16K Vocab)	11
7.2.3	Variant 3: WordPiece (16K Vocab)	12
7.3	Tokenizer Comparison Across All Variants	12
7.4	Tokenizer Evaluation Results	12
7.4.1	Telugu Byte-Level BPE (50K)	12
7.4.2	Telugu Unicode-Level BPE (50K)	12
7.4.3	Telugu WordPiece (50K)	13
7.4.4	Bhojpuri Byte-Level BPE (8K)	13
7.4.5	Bhojpuri Unicode-Level BPE (16K)	13
7.4.6	Bhojpuri WordPiece (16K)	13
7.4.7	Key Evaluation Metrics Summary	13
7.5	Tokenization Examples from Executed Notebooks	13
7.5.1	Telugu Byte-Level BPE (50K)	14
7.5.2	Telugu Unicode-Level BPE (50K)	14
7.5.3	Telugu WordPiece (50K)	14
7.5.4	Bhojpuri Byte-Level BPE (8K)	14
7.5.5	Bhojpuri Unicode-Level BPE (16K)	14
7.5.6	Bhojpuri WordPiece (16K)	14
8	Technical Implementation	14
8.1	Data Pipeline Architecture	14
8.2	File Structure	15
9	Quality Assurance	15
9.1	Validation Checklist	15
9.2	Reproducibility	16
10	Challenges & Solutions	16
10.1	Challenge 1: Bhojpuri Data Scarcity	16
10.2	Challenge 2: OCR Quality	16
10.3	Challenge 3: Scale Management	16
11	Dataset and Tokenizer Repository	16
11.1	Phase 1 Tokenizers	16
11.2	Dataset Contents	17

A Configuration Files	17
A.1 Telugu config.json Structure	17
A.2 Bhojpuri config.json Structure	17
B References	17

1 Phase 1 Data Report

This report provides a comprehensive data collection and preprocessing report for two monolingual Transformer language models:

- **Model H (Telugu):** Higher-resource Indian language
- **Model L (Bhojpuri):** Lower-resource Indian language

1.1 Data Collection Summary

1.1.1 Telugu (Model H)

- Total lines collected: 25,349,199
- Estimated tokens: 4,421,917,647 (4.4B)
- Web-scraped texts: 29,206 (89.7% pass rate from 32,549 raw)
- Manual corpus: 21,458,261 lines (te.txt)
- OCR augmentation: 3,766,703 lines
- Target: 500M tokens | Progress: 884.4%

1.1.2 Bhojpuri (Model L)

- Total lines collected: 776,801
- Estimated tokens: 136,606,695 (136.6M)
- Web-scraped texts: 3,234 (97.2% pass rate from 3,327 raw)
- HuggingFace Corpus: 386,000 documents (25M tokens)
- OCR augmentation: 511,599 lines
- Machine Translation (Hindi→Bhojpuri): 15,852 lines
- Target: 500M tokens | Progress: 27.3%

1.2 Data Processing Pipeline

1.2.1 Cleaning Process

- Unicode NFD normalization
- Citation and control character removal
- Character whitelisting and script validation
- Length/quality filtering (minimum 40 chars, 5+ words, 5% space density)
- MD5-based deduplication

1.2.2 Data Splits (80/10/10)

- Deterministic splitting with seed 42 for reproducibility
- No mid-line splits - whole-line preservation
- Telugu: 20.3M train / 2.5M val / 2.5M test
- Bhojpuri: 627.9K train / 78.1K val / 78.4K test

2 Data Collection Overview

2.1 Methodology

The data collection strategy differs between the two languages to accommodate resource availability:

2.1.1 Telugu (Model H) - Higher-Resource Language

- **Primary Source:** Existing corpus (te.txt) - 21.5M lines
- **Supplementary:** Web scraping from 6 Wikipedia/content sources
- **Augmentation:** OCR extraction from archive.org texts
- **Total:** 25.3M lines training data

2.1.2 Bhojpuri (Model L) - Lower-Resource Language

- **Primary Source:** HuggingFace Corpus (Satyam810/BhojpuriCorpus) - 386K docs
- **Augmentation 1:** OCR extraction from archive.org Hindi texts - 511.6K lines
- **Augmentation 2:** Web scraping from Hindi news sites (Devanagari proxy) - 3.2K lines
- **Augmentation 3:** Machine Translation (Hindi→Bhojpuri via NLLB-200) - 15.8K lines
- **Total:** 776.8K lines training data (27.3% toward 500M token target)

2.2 Data Sources

Language	Source	Contribution
Telugu	te.txt (existing)	21.5M lines (99.9%)
	te.wikipedia.org	3K articles
	te.wikibooks.org	500+ pages
	te.wikiquote.org	300+ quotes
	News portals	eenadu.net, greatandhra.com
	archive.org OCR	3.8M lines (augmentation)
Bhojpuri	HF Corpus	386K docs (~25M tokens)
	archive.org OCR	511.6K lines
	bho.wikipedia.org	~8.9K articles
	Hindi News	Via NLLB-200 MT
	Web Scraping	3.2K lines (Hindi proxy)
	GlottCC-V1	949 docs
	HPLT 3.0	~32.8K docs
	fish-food	262K rows

Table 1: Data sources for both languages

3 Data Cleaning Pipeline

3.1 7-Stage Cleaning Process

All data undergoes a rigorous 7-stage cleaning pipeline:

1. **Unicode Normalization:** Convert to NFD (Canonical Decomposition)
2. **Citation Removal:** Remove Wikipedia citations and references
3. **Control Character Removal:** Strip non-printable characters
4. **Character Whitelisting:** Retain only valid Unicode ranges
5. **Script Validation:** Language-specific script filtering
6. **Length/Quality Filters:** Minimum length, word count, density checks
7. **Deduplication:** MD5-based exact duplicate removal

3.2 Quality Metrics

Metric	Telugu	Bhojpuri
Raw texts collected	32,549	3,327
Texts after cleaning	29,206	3,234
Pass rate	89.7%	97.2%
Rejected texts	3,343	93

Table 2: Data cleaning results summary

3.2.1 Telugu Cleaning Analysis

- **Total Raw:** 32,549 web-scraped texts
- **Pass Rate:** 89.7% (29,206 cleaned)
- **Rejection Rate:** 10.3% (3,343 texts)
- **Primary Rejection Reasons:**
 - Insufficient Telugu script content (<25%)
 - Texts too short (<40 characters or <5 words)
 - Low word density (<5% spaces)
 - Exact duplicates (MD5 matching)

3.2.2 Bhojpuri Cleaning Analysis

- **Total Raw:** 3,327 web-scraped texts
- **Pass Rate:** 97.2% (3,234 cleaned)
- **Rejection Rate:** 2.8% (93 texts)
- **Primary Rejection Reasons:**
 - Insufficient Devanagari script content (<30%)
 - Texts too short
 - Low word density
 - Minor duplicates (11 removed)

4 Data Statistics & Visualizations

4.1 Collection Metrics Comparison




Figure 1: **Left:** Total lines per language. Telugu has 32.5x more training data than Bhojpuri. **Right:** Estimated tokens. Telugu already exceeds 500M target; Bhojpuri at 27.3% of target.

4.2 Cleaning Pipeline Results




Figure 2: Data cleaning pipeline effectiveness. Both languages show strong pass rates. Telugu: 89.7% pass rate with large-scale data. Bhojpuri: 97.2% pass rate with targeted scraping.

4.3 Train/Validation/Test Split Distribution




Figure 3: Deterministic 80/10/10 train/validation/test splits using seed 42. Both languages follow identical split ratio for reproducibility.

4.4 Data Source Composition (Bhojpuri)




Figure 4: Bhojpuri data comes from diverse sources. HuggingFace Corpus (386K docs) is primary source, with OCR augmentation (511.6K lines) providing significant additional content.

5 Train/Validation/Test Splits

5.1 Split Strategy

- **Method:** 80/10/10 deterministic split with seed 42
- **Reproducibility:** Same seed ensures identical splits across runs
- **No Data Leakage:** Splits applied at line level (no mid-sentence splits)

- **Whole-Line Preservation:** Each line remains intact for language preservation

5.2 Telugu Splits

Split	Lines	Percentage
Train	20,281,561	80.0%
Validation	2,534,020	10.0%
Test	2,533,618	10.0%
Total	25,349,199	100.0%

Table 3: Telugu (Model H) data split breakdown

5.3 Bhojpuri Splits

Split	Lines	Percentage
Train	627,925	80.0%
Validation	78,137	10.0%
Test	78,391	10.0%
Total	776,801	100.0%

Table 4: Bhojpuri (Model L) data split breakdown

6 Detailed Statistics

6.1 Telugu (Model H) Statistics

Metric	Value
Language	Telugu
Model Size	Higher-resource
<i>Data Collection</i>	
Web scraping - Raw texts	32,549
Web scraping - Cleaned texts	29,206
Web scraping - Size	27.6 MB
Manual corpus (te.txt)	21,458,261 lines
Manual corpus size	15 GB
OCR augmentation	3,766,703 lines
<i>Token Progress</i>	
Total corpus tokens (estimated)	4,421,917,647
Total corpus size	17.7 GB
Target tokens	500,000,000
Progress	884.4% (EXCEEDED)

Table 5: Comprehensive Telugu statistics

6.2 Bhojpuri (Model L) Statistics

Metric	Value
Language	Bhojpuri
Model Size	Lower-resource
<i>Data Collection</i>	
Web scraping - Raw texts	3,327
Web scraping - Cleaned texts	3,234
Web scraping - Pass rate	97.2%
HuggingFace Corpus	386,000 docs
Estimated HF tokens	25,000,000
OCR augmentation (archive.org)	511,599 lines
OCR augmentation size	19.7 MB
Machine Translation (Hindi→Bhojpuri)	15,852 lines
Translation source	Hindi news + Wikipedia
Translation method	NLLB-200 distilled-600M
<i>Token Progress</i>	
Total corpus tokens (estimated)	136,606,695
Total corpus size	546.4 MB
Target tokens	500,000,000
Progress	27.3%
Data Ready for Phase 2	YES

Table 6: Comprehensive Bhojpuri statistics

7 Tokenizer Training

Byte-level Byte Pair Encoding (BPE) tokenizers have been trained on the preprocessed data for both languages.

7.1 Telugu Tokenizer (Model H) - Three Variants

Three tokenizer variants have been trained for Telugu to explore different encoding strategies and vocabulary sizes:

7.1.1 Variant 1: Byte-Level BPE (50K Vocab)

Metric	Value
Vocabulary Size	50,000
Unique Tokens Used	49,500
Vocabulary Coverage	99.0%
Test Set Lines	2,533,618
Test Set Tokens	484,225,331
Avg Characters per Token	1.36
Unknown (UNK) Rate	0.0%

Table 7: Telugu byte-level tokenizer (50K vocab)

7.1.2 Variant 2: Unicode-Level BPE (50K Vocab)

Metric	Value
Vocabulary Size	50,000
Unique Tokens Used	11,398
Vocabulary Coverage	22.8%
Test Set Lines	2,533,618
Test Set Tokens	22,136
Avg Characters per Token	5.9283
Unknown (UNK) Rate	0.0%

Table 8: Telugu unicode-level tokenizer (50K vocab)

7.1.3 Variant 3: WordPiece (50K Vocab)

Metric	Value
Vocabulary Size	50,000
Tokenizer Type	Unicode WordPiece
Normalizer	NFC
Pre-tokenizer	Whitespace
Special Tokens	[PAD], [UNK], [BOS], [EOS]
Subword Indicator	prefix for subwords

Table 9: Telugu WordPiece tokenizer (50K vocab)

7.2 Bhojpuri Tokenizer (Model L) - Three Variants

Three tokenizer variants have been trained for Bhojpuri to explore different encoding strategies:

7.2.1 Variant 1: Byte-Level BPE (8K Vocab)

Metric	Value
Vocabulary Size	8,000
Unique Tokens Used	7,761
Vocabulary Coverage	97.01%
Test Set Lines	77,840
Test Set Tokens	13,613,342
Avg Characters per Token	1.54
Unknown (UNK) Rate	0.0%

Table 10: Bhojpuri byte-level tokenizer (8K vocab)

7.2.2 Variant 2: Unicode-Level BPE (16K Vocab)

Metric	Value
Vocabulary Size	16,000
Unique Tokens Used	15,643
Vocabulary Coverage	97.77%
Test Set Lines	77,840
Test Set Tokens	6,072,593
Avg Characters per Token	3.44
Unknown (UNK) Rate	0.0%

Table 11: Bhojpuri unicode-level tokenizer (16K vocab)

7.2.3 Variant 3: WordPiece (16K Vocab)

Metric	Value
Vocabulary Size	16,000
Tokenizer Type	WordPiece
Random Seed	42 (reproducible)
Test Set Tokens	5,930,926
Unknown (UNK) Rate	0.0001%
File Size	453 KB
Training Status	Completed

Table 12: Bhojpuri WordPiece tokenizer (16K vocab)

7.3 Tokenizer Comparison Across All Variants

Language	Variant	Vocab	Coverage	UNK Rate
Telugu	Byte-level (BPE)	50,000	3.9%	0.0%
	Unicode-level (BPE)	50,000	22.8%	0.0%
	WordPiece	50,000	19.1%	0.0%
Bhojpuri	Byte-level (BPE)	8,000	97.01%	0.0%
	Unicode-level (BPE)	16,000	97.77%	0.0%
	WordPiece	16,000	Excellent	0.0001%

Table 13: Complete tokenizer variants summary

7.4 Tokenizer Evaluation Results

Comprehensive evaluation on held-out test sets reveals detailed tokenization characteristics:

7.4.1 Telugu Byte-Level BPE (50K)

- **Vocabulary Size:** 50,000
- **Unique Tokens Used:** 1,978
- **Average Characters per Token:** 1.3327
- **Unknown Token Rate:** 0.0000%
- **Top Token:** ____ (10.12% of tokens)

7.4.2 Telugu Unicode-Level BPE (50K)

- **Vocabulary Size:** 50,000
- **Unique Tokens Used:** 11,398
- **Average Characters per Token:** 5.9283
- **Unknown Token Rate:** 0.0000%
- **Top Tokens:** Functional words (లో, ఈ, కూడా)

7.4.3 Telugu WordPiece (50K)

- **Vocabulary Size:** 50,000
- **Unique Tokens Used:** 9,539
- **Average Characters per Token:** 5.8353
- **Unknown Token Rate:** 0.0000%
- **Subword Strategy:** Explicit morphological decomposition with # prefix

7.4.4 Bhojpuri Byte-Level BPE (8K)

- **Vocabulary Size:** 8,000
- **Average Characters per Token:** 1.54
- **Coverage:** 97.01%
- **Unknown Token Rate:** 0.0%

7.4.5 Bhojpuri Unicode-Level BPE (16K)

- **Vocabulary Size:** 16,000
- **Average Characters per Token:** 3.44
- **Coverage:** 97.77%
- **Unknown Token Rate:** 0.0%
- **Top Tokens:** आ 9.7%, ु 7.1%, े 5.8%

7.4.6 Bhojpuri WordPiece (16K)

- **Vocabulary Size:** 16,000
- **Test Tokens Evaluated:** 5,930,926
- **Unknown Token Rate:** 0.0001%
- **Top Tokens:** के, l, comma

7.4.7 Key Evaluation Metrics Summary

Tokenizer	Avg Chars/Token	UNK Rate	Tokens Used
Telugu Byte-level	1.3327	0.0000%	1,978
Telugu Unicode-level	5.9283	0.0000%	11,398
Telugu WordPiece	5.8353	0.0000%	9,539
Bhojpuri Byte-level	1.54	0.0%	Unknown
Bhojpuri Unicode-level	3.44	0.0%	Unknown
Bhojpuri WordPiece	3-5 (est.)	0.0001%	5.9M test

Table 14: Tokenizer evaluation metrics summary

7.5 Tokenization Examples from Executed Notebooks

One example per tokenizer from test set evaluation:

7.5.1 Telugu Byte-Level BPE (50K)

Input Text: నందమూరి తారక రామారావు 50 సంవత్సరాలపైగా తెలుగు సినిమా రంగంలో కథా నాయకునిగా రాణించాడు

Tokenized Output: Tokens: ['న', 'ందమూరి', ' ', 'తారక', ' ', 'రామారావు', ' ', '50', ' ', 'సం', 'వత్సరాల', 'పైగా', ' ', 'తెలుగు', ' ', 'సినిమా', ' ', 'రంగంలో', ' ', 'కథా', '...']

Total tokens: 91

7.5.2 Telugu Unicode-Level BPE (50K)

Input Text: నందమూరి తారక రామారావు 50 సంవత్సరాలపైగా తెలుగు సినిమా రంగంలో కథా నాయకునిగా రాణించాడు

Tokenized Output: Tokens: ['నందమూరి', 'తారక రామారావు', '50', 'సంవత్సరాల', 'పైగా', 'తెలుగు సినిమా', 'రంగంలో', 'కథా', 'నాయకు', 'నిగా', '...']

Total tokens: 21

7.5.3 Telugu WordPiece (50K)

Input Text: నందమూరి తారక రామారావు 50 సంవత్సరాలపైగా తెలుగు సినిమా రంగంలో కథా నాయకునిగా రాణించాడు

Tokenized Output: Tokens: ['నందమూరి', 'తారక', 'రామారావు', '50', 'సంవత్సరాల', 'పైగా', 'తెలుగు', 'సినిమా', 'రంగంలో', 'కథా', 'న', 'గా', '...']

Total tokens: 21

7.5.4 Bhojpuri Byte-Level BPE (8K)

Input Text: बहुत बढ़िया काम है यह।

Tokenized Output: Tokens: ['बहु', 'त', ' ', 'बढ़', 'ि', 'िया', ' ', 'क', 'ाम', ' ', '...']

Total tokens: varies by script

7.5.5 Bhojpuri Unicode-Level BPE (16K)

Input Text: भोजपुरी भारत के बिहार में बोली जाती है।

Tokenized Output: Tokens: ['भोजपुरी', ' ', 'भारत', ' ', 'के', ' ', 'बिहार', ' ', 'में', ' ', '...']

Total tokens: varies by script

7.5.6 Bhojpuri WordPiece (16K)

Input Text: भोजपुरी भारत के बिहार में बोली जाती है।

Tokenized Output: Token IDs: [662, 748, 382, 1076, 388, 3827, 1471, 390, 218]

Tokens: [भोजपुरी, भारत, के, बिहार, में, बोली, जाती, है, ।]

Total tokens: 9

8 Technical Implementation

8.1 Data Pipeline Architecture

- **Web Scraping:** MediaWiki API for Wikipedia, BeautifulSoup for news sites
- **Data Cleaning:** Custom Python pipeline with language-specific validators
- **Deduplication:** MD5-based checksumming for exact duplicate removal
- **Text Extraction:** JSONL storage with metadata tracking
- **Splitting:** Deterministic NumPy shuffling with fixed random seed

8.2 File Structure

Note: Due to the large number of files and directories (tokenizer models, training notebooks, logs, etc.), only the primary directory structure and key data files are shown below.

```
telugu/
  configs/                # Configuration files & parameters
  data_collect/           # Data collection scripts
    scraper.py
    data_cleaner.py
    pipeline.py
  data/                   # Preprocessed dataset (17.4 GB total)
    config.json           # Complete statistics & metadata
    scrape_state.json     # Scraping checkpoint
    train/                # 20.3M lines (14 GB, 80%)
    val/                  # 2.5M lines (1.7 GB, 10%)
    test/                 # 2.5M lines (1.7 GB, 10%)
  model/                  # Model architecture & checkpoints
  tokenizer/              # Trained tokenizers (6 variants)
    full_byte_level/
    full_unicode_level/
    full_wordPiece_level/
  train/                  # Training scripts
  eval/                   # Evaluation scripts

bhojpuri/
  configs/                # Configuration files & parameters
  data_collect/           # Data collection scripts
    scraper.py
    data_cleaner.py
    pipeline.py
  data/                   # Preprocessed dataset (518 MB total)
    config.json           # Complete statistics & metadata
    scrape_state.json     # Scraping checkpoint
    train/                # 627.9K lines (414 MB, 80%)
    val/                  # 78.1K lines (52 MB, 10%)
    test/                 # 78.4K lines (52 MB, 10%)
  model/                  # Model architecture & checkpoints
  tokenizer/              # Trained tokenizers (3 variants)
    full_byte_level/
    full_unicode_level/
    full_wordPiece_level/
  train/                  # Training scripts
  eval/                   # Evaluation scripts
```

9 Quality Assurance

9.1 Validation Checklist

- ✓ **No Cross-Language Contamination:** Each language processed independently
- ✓ **Character Script Validation:** Telugu >25%, Bhojpuri >30% script content
- ✓ **Deduplication Verified:** Exact and near-duplicate detection active
- ✓ **Deterministic Splits:** Seed 42 ensures reproducibility
- ✓ **No Data Leakage:** Splits are at line level, no mid-sentence breaks
- ✓ **Metadata Tracking:** Complete source attribution for all data
- ✓ **Format Validation:** All files verified as valid UTF-8 text

9.2 Reproducibility

All data processing is fully reproducible:

- **Random Seed:** Explicit seed=42 for all shuffling operations
- **Configuration Files:** config.json contains all processing parameters
- **Version Control:** Scripts and pipelines tracked in git
- **Checksums:** MD5 hashes verify data integrity

10 Challenges & Solutions

10.1 Challenge 1: Bhojpuri Data Scarcity

- **Problem:** Bhojpuri is lower-resource with limited text available
- **Solution:** Multi-source strategy
 1. Primary: HuggingFace Corpus (386K docs)
 2. Augmentation 1: OCR extraction from Hindi texts
 3. Augmentation 2: Machine translation (Hindi→Bhojpuri)
 4. Augmentation 3: Wikipedia dumps + web scraping
- **Result:** 776.8K lines collected (sufficient for Phase 2-3)

10.2 Challenge 2: OCR Quality

- **Problem:** Archive.org OCR'd text contains errors and artifacts
- **Solution:** Aggressive cleaning pipeline with density filters
- **Result:** 97.2% pass rate even on noisy OCR data

10.3 Challenge 3: Scale Management

- **Problem:** Telugu manual corpus is 15GB (21.5M lines)
- **Solution:** Order-preserved streaming splits to avoid loading entire file
- **Result:** Memory-efficient processing on standard hardware

11 Dataset and Tokenizer Repository

Phase 1 Data The complete Phase 1 dataset is available on Kaggle:

<https://www.kaggle.com/datasets/kpsvlnsiddardha/lma-slm>

11.1 Phase 1 Tokenizers

All 6 trained tokenizers (Telugu + Bhojpuri variants) are available on Kaggle:

<https://www.kaggle.com/datasets/kpsvlnsiddardha/lma-tokenizers>

11.2 Dataset Contents

- **Telugu Training Data:** 20.3M lines (80% of 25.3M total)
- **Telugu Validation Data:** 2.5M lines (10% of 25.3M total)
- **Telugu Test Data:** 2.5M lines (10% of 25.3M total)
- **Bhojpuri Training Data:** 627.9K lines (80% of 776.8K total)
- **Bhojpuri Validation Data:** 78.1K lines (10% of 776.8K total)
- **Bhojpuri Test Data:** 78.4K lines (10% of 776.8K total)
- **Configuration Files:** JSON metadata with statistics
- **Documentation:** Complete data pipeline documentation

A Configuration Files

A.1 Telugu config.json Structure

```
{
  "language": "Telugu",
  "data_sources": [...],
  "scraping": {
    "raw_texts": 32549,
    "texts_cleaned": 29206,
    "pass_rate_percent": 89.7
  },
  "splits": {
    "train": {"lines": 20281561, "percentage": 0.80},
    "val": {"lines": 2534020, "percentage": 0.10},
    "test": {"lines": 2533618, "percentage": 0.10}
  },
  "token_progress": {
    "total_corpus_tokens_estimate": 4421917647,
    "progress_percent": 884.3835
  }
}
```

A.2 Bhojpuri config.json Structure

Similar structure with Bhojpuri-specific sources and metrics (see main config file).

B References

- **Kaggle Dataset:** <https://www.kaggle.com/datasets/kspsvlnsiddardha/lma-slm>
- **Project Repository:** /media/ubuntu/Personal/IIIT Hyderabad/Semester 3/LMA/Mini_Project/
- **README:** README.md
- **Language Documentation:** bhojpuri/BHOJPURI.md, telugu/TELUGU.md
- **Configuration Files:** telugu/data/config.json, bhojpuri/data/config.json